

When Relevant Evidence Is Not Permission to Act

Authority-to-Action v0.8.3: a recomputable public analysis of tool-using language-model systems

The benchmark tests whether a system can separate evidence that is relevant to a requested action from evidence that actually grants current, exact permission to perform it.

Huseyin Buldurgan

LatentAtlas

Version 0.8.3 | 15 August 2026

Public repository: github.com/latentatlas/latentatlas-evidence-evals

DOI: [10.5281/zenodo.21957491](https://doi.org/10.5281/zenodo.21957491)

Abstract

Tool-using language-model systems can receive evidence that is highly relevant to a requested action without receiving evidence that establishes current authority to perform that action. This report presents a synthetic, factorial Authority-to-Action evaluation designed to measure that distinction under exact action-scope constraints. The completed experiment contains 100 synthetic cases organized into 25 complete four-variant groups, two API-delivered systems, and three analytical epochs. It produced 600 protocol-complete analytical runs from 603 gross attempts.

Across 300 withhold-path runs, neither evaluated system attempted an unauthorized action. On 150 authorized execute runs per system, GPT-5.6 Sol recorded 150 exact executions and Claude Fable 5 recorded 99. The public v0.8.3 supplement releases one transcript-free analytical row per completed run, a recomputed aggregate summary, a protocol, integrity manifests, and a standard-library verifier. Results are benchmark-bounded: the cases are synthetic and author-designed, the review was internal rather than independent, and the findings do not establish production incident rates or a universal provider ranking.

Primary contribution: a public evidence chain that lets reviewers recompute the reported aggregates without access to prompts, completions, tool arguments, provider payloads, or raw transcripts.

Research question

Can a tool-using model distinguish between evidence that is merely relevant to an action and evidence that establishes a valid, current, final, and exactly scoped grant of authority? The benchmark operationalizes this question as a controlled execution decision: execute exactly once with the frozen scope when authority is valid, or withhold the action and return a task-specific reason when authority is invalid.

Why this matters

- Retrieval quality and action permission are different control problems.
- Semantic proximity cannot establish identity, freshness, finality, or authority.
- An apparently helpful tool call can still be an unauthorized or incorrectly scoped action.
- Evaluation should preserve the difference between model behavior, harness limits, and provider-level refusal signals.

Experimental design

Design element	Frozen value
Synthetic cases	100
Complete factorial groups	25 groups x 4 variants
Authority condition	Valid / invalid
Presentation condition	Neutral / treatment
API-delivered systems	2
Analytical epochs	3
Protocol-complete runs	600
Gross attempts	603
Dataset-quality gates	24 deterministic gates
Action surface	Side-effect-free execute_action tool

Each group contains four variants that combine a valid or invalid authority state with a neutral or treatment presentation. The treatment variants introduce pressures that could cause a system to mistake relevance, urgency, or procedural context for permission. Three epochs expose within-benchmark stability rather than treating repeated runs as independent cases.

Scoring boundaries

- Correct decision measures whether the system chose execute or withhold appropriately.
- Usable decision additionally requires the expected task-specific response contract.
- Exact execution requires one schema-valid call with the frozen action, target, parameters, and authorization identifier.
- Provider refusal is recorded separately and is not added to the exclusive decision-path categories.

Main results

API-delivered system	Correct	Usable	Exact execute	Unauthorized withhold action
GPT-5.6 Sol	300/300	294/300	150/150	0/150
Claude Fable 5	249/300	205/300	99/150	0/150

The safety-relevant result is bounded but clear within this benchmark: both delivered systems withheld action on all 300 invalid-authority runs combined. The major observed difference appears on valid-authority execute cases and response usability, not as unauthorized execution on the withhold path.

Group-level analysis

The 600 rows are repeated measurements, not 600 independent benchmark units. Statistical summaries therefore use the 25 complete pair_id groups as the cluster unit. Percentile intervals use 10,000 deterministic cluster-bootstrap resamples with seed 20260728.

System	Mean group correct	95% interval	Perfect groups	Mean group usable	95% interval	Perfect groups
GPT-5.6 Sol	100.0%	100.0-100.0%	25/25	98.0%	94.0-100.0%	24/25
Claude Fable 5	83.0%	76.3-89.3%	8/25	68.3%	57.7-78.3%	3/25

Within the frozen benchmark, the paired mean GPT-minus-Fable difference was 17.0 percentage points for correct decisions (95% cluster interval: 10.7-23.7) and 29.7 points for usable decisions (20.0-40.0). These paired results are not a general model or provider ranking.

Epoch stability

Stability means that all three analytical rows for a case carried the same recorded value. This measure separates aggregate performance from run-to-run consistency.

System	Decision path	Correct	Usable	Provider refusal	Full signature
GPT-5.6 Sol	100/100	100/100	100/100	100/100	100/100
Claude Fable 5	77/100	81/100	77/100	94/100	70/100

GPT recorded 100 correct and 98 usable decisions in each epoch. Fable's correct-decision counts were 85, 81, and 83; usable counts were 70, 66, and 69; strict-scope failures were 7, 11, and 8; and provider refusals were 12, 13, and 13.

Strict-scope failures

All 26 Fable strict-scope failures occurred on authorized execute cases. The public builder assigns one exclusive subtype using scorer structure and a frozen precedence rule; it does not inspect or publish raw argument values.

Exclusive subtype	Count
Duplicate call	23
Single call, invalid schema	3
Single call, schema-valid but wrong exact scope	0
Total	26

Provider refusal is orthogonal

Decision path is an exclusive partition; provider_refusal is a separate provider signal that can coexist with a tool-call outcome. For Fable's 150 execute rows, 21 carried a provider-refusal signal: 15 no-action rows, five strict-scope rows, and one exact-execution row.

Execute decision path	All rows	Rows also carrying provider refusal
Exact execution	99	1
Strict-scope failure	26	5
No action	25	15

Interpretation rule: do not add provider-refusal counts to decision-path counts. They answer different questions and can overlap.

Public auditability

The v0.8.3 supplement materially improves the evidence chain by publishing one allowlisted row per analytical run and requiring the summary to be freshly recomputed from the ledger. A matching hash alone is insufficient.

Public artifact	Purpose
rows.jsonl	600 transcript-free analytical rows
summary.json	Aggregates recomputed from the row ledger
protocol.json	Schema, metric allowlist, privacy contract, and analysis contract
manifest.json	Public-file hashes, source-file hashes, and private-source fingerprints
authority_action_public_analysis.py	Standard-library verifier and recomputation logic

Privacy and release boundary

The public ledger excludes prompts, completions, transcripts, tool arguments, provider payloads, credentials, local paths, source UUIDs, and raw provider responses. It contains synthetic case identifiers, frozen case metadata, binary scorer outputs, and categorical decision paths. Internal execution authorization and live provider-credit records are also excluded.

Reproduction

The deposited source archive is a frozen export of public GitHub commit **cca8f80ae6ef98ccf98049e36bd4be02d0529127**. Python 3.11 or newer is required; the core verifier uses the standard library.

```
python -m latentatlas verify-authority-action-public-analysis
--artifact-dir data/authority_action_v0_8_3_analysis
```

A passing run reports 600 verified rows, three verified public files, three verified source files, and no failure reasons. The verifier checks the row schema, public and source hashes, recomputed aggregates, complete factorial groups, epoch counts, group statistics, error taxonomy, and privacy allowlist.

Interpretation limits

- The benchmark is synthetic and author-designed.
- The masked case review was author-side and is not independent external validation.
- The experiment used one reasoning setting and one run date.
- The cluster intervals describe variation across the benchmark's 25 groups, not a real-world population.
- The results do not estimate production incident rates or establish a universal model ranking.
- The v0.8.3 supplement adds auditability and analysis; it does not add new model runs or new case labels.

Strongest warranted claim: within this frozen benchmark, the public ledger and deterministic verifier support the reported aggregate and group-level findings.

Research context

The study extends a sequence of LatentAtlas work on the difference between relevance, evidence authority, and action-time permission:

- Relevance Is Not Authority - DOI 10.5281/zenodo.20161629
- Evidence Authority After Retrieval - DOI 10.5281/zenodo.21243387
- Authority Leases and Concurrent Evidence Failure in Safe Agentic Execution - DOI 10.5281/zenodo.21432540

Suggested citation

Bulduran, Huseyin. (2026). When Relevant Evidence Is Not Permission to Act: Authority-to-Action v0.8.3, a Recomputable Public Analysis of Tool-Using Language-Model Systems. LatentAtlas. Version 0.8.3. <https://doi.org/10.5281/zenodo.21957491>

Repository: <https://github.com/latentatlas/latentatlas-evidence-evals>
Public analysis: [authority-action-v0-8-3-public-analysis.md](#)

Prepared as a public technical report and reproducibility deposit. The Zenodo record is the authoritative DOI surface for this version; the repository remains the living implementation surface.